



Black and White to RGB Color

ECS 171 Group 3

Fall 2025

Instructor Ilias Tagkopoulos



AUTO

ISO 90

ISO 100

ISO 200

ISO 400

ISO 600

Table of contents

01

Introduction

Why automated image
colorization is helpful

02

Background & Math

The core mathematical
principles of our approach

03

Methods

How we designed our
model

04

Results & Demo

Evaluating our model
performance + Demo

05

Discussion

Interpreting our results

06

Conclusion

Limitations, takeaways, and
future improvements



01 Introduction

AUTO

ISO 90

ISO 100

ISO 200

ISO 400

ISO 600

The Manual Approach

Background

- Older photos and films exist only in black and white
- Colorization done by hand by referencing knowledge of natural colors and historical trends



Problem

- Manual restoration methods are time-consuming and expensive
 - Limits accessibility and scalability

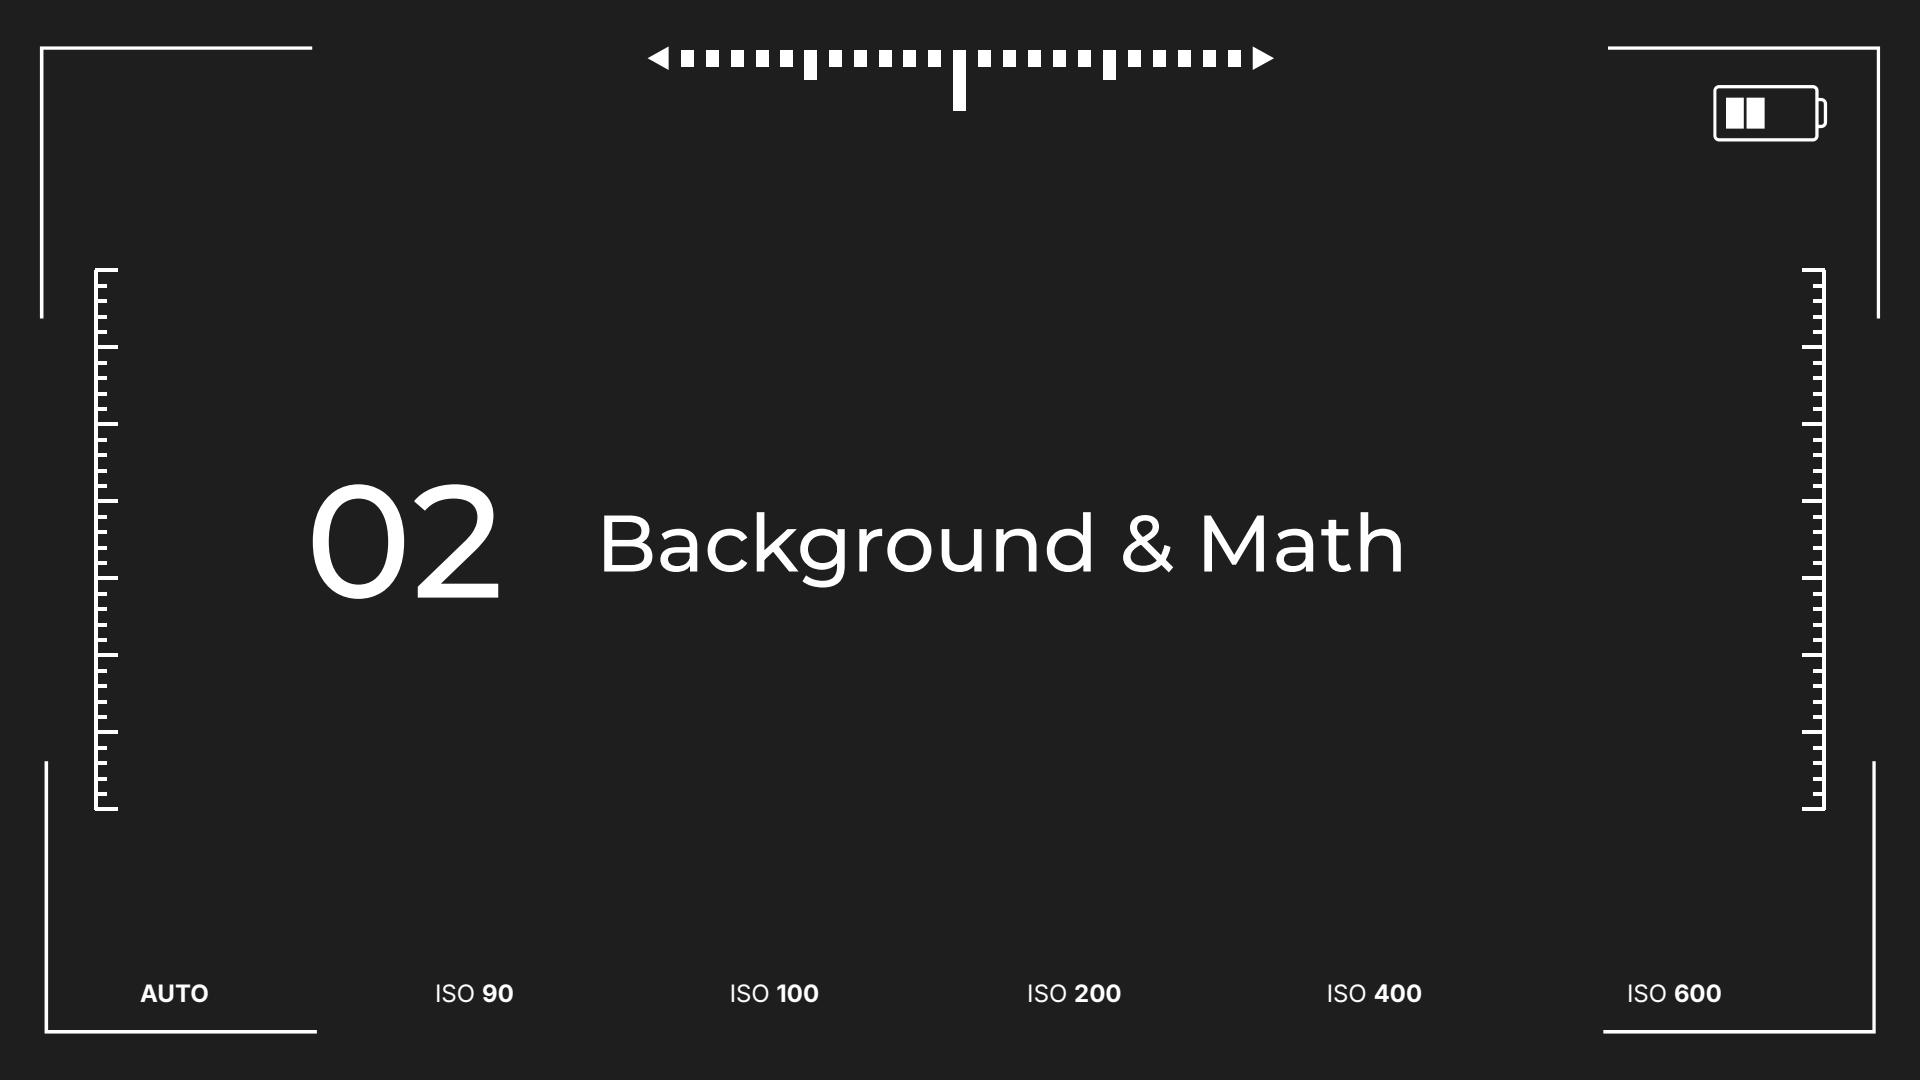




Our Idea

Use concepts learned in class with
computer vision ideas to make a
Convolution Neural Network (CNN)





02 Background & Math

AUTO

ISO 90

ISO 100

ISO 200

ISO 400

ISO 600



.....|.....

How do we represent color?





The Math: Grayscale

$$Y \in \mathbb{R}^{H \times W}$$

Height (pixels) *Width (pixels)*

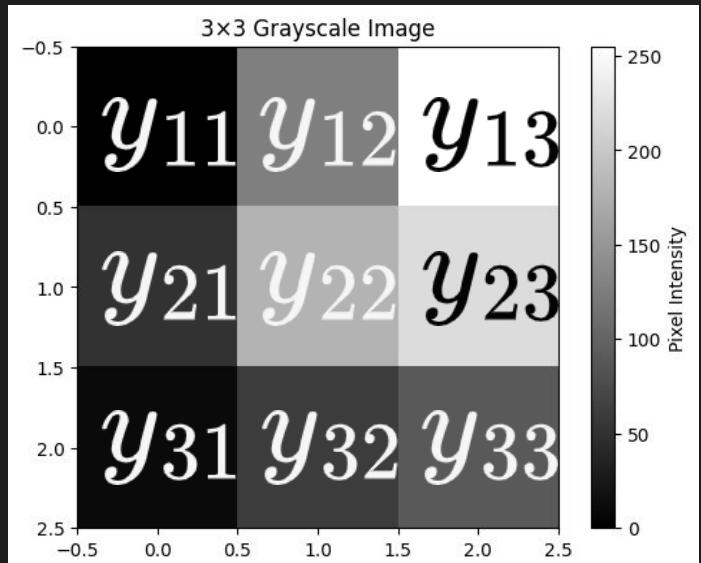
**2D Grayscale
Image (matrix)**

Each pixel in the image is stored as a scalar in the matrix, which represents the brightness of the pixel. (ranges from 0: black → 255: white)

$$y_{ij} \in \{0, 1, \dots, 255\}$$

**Intensity/Brightness
of Pixel (scalar)**

$$Y^{3 \times 3} = \begin{bmatrix} 0 & 127 & 255 \\ 50 & 180 & 220 \\ 10 & 60 & 90 \end{bmatrix}$$





The Math: Colored

$$V \in \mathbb{R}^{H \times W \times 3}$$

Height (pixels) *Width (pixels)* *RGB Color Channels*

3D Colored Image (tensor)

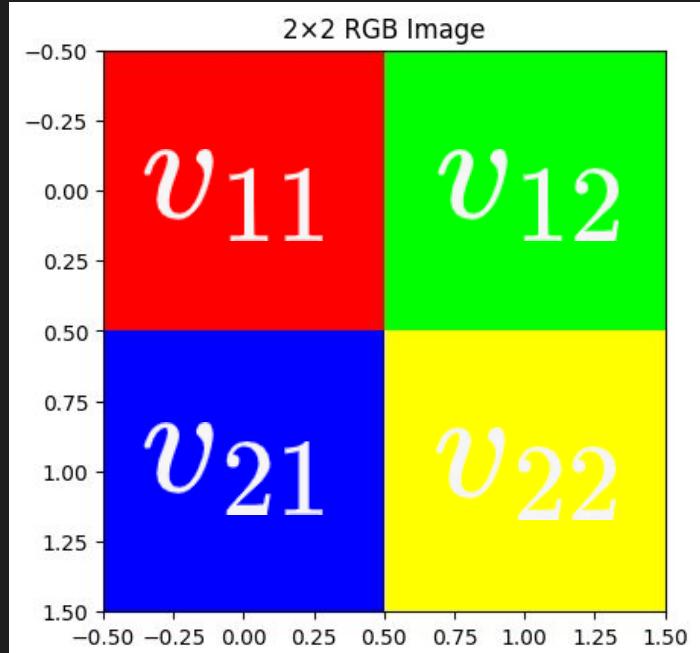
Each pixel is stored as a 3D vector, where each component represents the intensity of the **red**, **green**, and **blue** color channels, respectively.

$$\mathbf{v}_{ij} = \begin{bmatrix} R_{ij} \\ G_{ij} \\ B_{ij} \end{bmatrix} \in \mathbb{R}^3 \in \{0, 1, \dots, 255\}$$

Color Intensity of Pixel (scalar)

3D pixel vector

$$V^{2 \times 2 \times 3} = \begin{bmatrix} [(255, 0, 0)] & [(0, 255, 0)] \\ [(0, 0, 255)] & [(255, 255, 0)] \end{bmatrix}$$





From **RGB** to Grayscale

We want:

$$f(I_{\text{gray}}) \rightarrow I_{\text{color}}$$

$$f(\mathbb{R}^{H \times W}) \rightarrow \mathbb{R}^{H \times W \times 3}$$

2D Grayscale Matrix

3D Colored Tensor

Convert grayscale to **RGB**

$$Y \in \mathbb{R}^{H \times W} \rightarrow V \in \mathbb{R}^{H \times W \times 3}$$
$$y_{ij} \in \mathbb{R} \quad \quad \quad \mathbf{v}_{ij} \in \mathbb{R}^3$$

2D Grayscale Matrix

3D **RGB Tensor**

We can project the grayscale part y_{ij} as a linear combination of the **RGB** components of the colored pixel \mathbf{v}_{ij} .

W is a constant weight vector that specifies the contribution of each RGB channel to pixel brightness.

$$\mathbf{w} = [0.299 \ 0.587 \ 0.114]^\top$$

given standard BT.601 weights

$$\begin{aligned} y_{ij} &= \mathbf{w}^\top \mathbf{v}_{ij} \\ &= [w_R \ w_G \ w_B] \begin{bmatrix} R_{ij} \\ G_{ij} \\ B_{ij} \end{bmatrix} \\ &= w_R R_{ij} + w_G G_{ij} + w_B B_{ij} \end{aligned}$$



.....|.....

How do we predict the
colored vector?





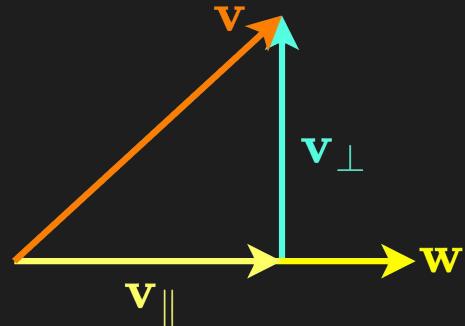
From Grayscale to RGB

$$\begin{aligned}\mathbf{v} &= \text{proj}_{\mathbf{w}}(\mathbf{v}) + (\mathbf{v} - \text{proj}_{\mathbf{w}}(\mathbf{v})) \\ &= \underbrace{\frac{\mathbf{w}^\top \mathbf{v}}{\mathbf{w}^\top \mathbf{w}} \mathbf{w}}_{\text{parallel to } \mathbf{w}} + \underbrace{\left(\mathbf{v} - \frac{\mathbf{w}^\top \mathbf{v}}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} \right)}_{\text{orthogonal to } \mathbf{w}} \\ &= \underbrace{\frac{y}{\mathbf{w}^\top \mathbf{w}} \mathbf{w}}_{\text{brightness part of } \mathbf{v}} + \underbrace{\left(\mathbf{v} - \frac{y}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} \right)}_{\text{the missing part of } \mathbf{v}} \quad \text{as } y = \mathbf{w}^\top \mathbf{v}\end{aligned}$$

$$\begin{aligned}&= y \underbrace{\left(\frac{1}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} \right)}_{\substack{\text{luminosity} \\ \text{scalar} \\ \in \{0, 1, \dots, 255\}}} + \underbrace{\left(\mathbf{v} - y \frac{1}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} \right)}_{\substack{\text{direction of } \mathbf{w} \\ \text{color component of } \mathbf{v} \\ \text{without luminosity} \\ (\text{chroma})}} \\ &= \mathbf{v}_{\parallel} + \mathbf{v}_{\perp}\end{aligned}$$

We want:

Solve for RGB vector \mathbf{v} given $y = \mathbf{w}^\top \mathbf{v}$.



RGB vector \mathbf{v} can be decomposed into a luminosity component \mathbf{v}_{\parallel} (brightness) and a chroma component \mathbf{v}_{\perp} (color without brightness).

This gives us the following constraint:

$$\mathbf{w}^\top (\mathbf{v} - \text{proj}_{\mathbf{w}}(\mathbf{v})) = \mathbf{w}^\top \mathbf{v}_{\perp} = 0$$



From Grayscale to **RGB**

We have the constraint: $\mathbf{W}^\top \mathbf{v}_\perp = 0$

Expanding this, we get:

$$w_R v_R + w_G v_G + w_B v_B = 0$$

Then, we can solve for one variable, say v_B :

Rewriting \mathbf{v}_\perp with this:

$$\mathbf{v}_\perp = \begin{bmatrix} v_R \\ v_G \\ -\frac{w_R v_R + w_G v_G}{w_B} \end{bmatrix} = v_R \begin{bmatrix} 1 \\ 0 \\ -\frac{w_R}{w_B} \end{bmatrix} + v_G \begin{bmatrix} 0 \\ 1 \\ -\frac{w_G}{w_B} \end{bmatrix}$$

$$\mathbf{w} = \begin{bmatrix} w_R \\ w_G \\ w_B \end{bmatrix} \quad \mathbf{v}_\perp = \begin{bmatrix} v_R \\ v_G \\ v_B \end{bmatrix}$$

$$v_B = -\frac{w_R v_R + w_G v_G}{w_B}$$

v_R and v_G are free parameters, while v_B 's value is constrained by orthogonality to \mathbf{W} .

The constraint defines a 2D subspace of \mathbb{R}^3 , so \mathbf{v}_\perp can be expressed as a linear combination of two linearly independent basis vectors.



From Grayscale to **RGB**

Choose any two linearly independent vectors:

$$\mathbf{u}_1 = \begin{bmatrix} u_{1R} \\ u_{1G} \\ u_{1B} \end{bmatrix}, \mathbf{u}_2 = \begin{bmatrix} u_{2R} \\ u_{2G} \\ u_{2B} \end{bmatrix} \in \mathbb{R}^3 \text{ such that } \mathbf{w}^\top \mathbf{u}_1 = \mathbf{w}^\top \mathbf{u}_2 = 0$$

Orthogonality Constraint

Then any possible chroma vector \mathbf{v}_\perp can be represented as a linear combination of them.

$$\mathbf{v}_\perp = \alpha \mathbf{u}_1 + \beta \mathbf{u}_2, \quad \alpha, \beta \in \mathbb{R}$$

So we can finally represent the RGB vector \mathbf{v} as:-

$$\mathbf{v} = \frac{y_{ij}}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} + \underbrace{\alpha \mathbf{u}_1 + \beta \mathbf{u}_2}_{\mathbf{v}_\perp}$$

^ Luminance

^ Chroma

We want:

Solve for RGB vector \mathbf{v}
given $y = \mathbf{w}^\top \mathbf{v}$.





03 Data and Methods

AUTO

ISO 90

ISO 100

ISO 200

ISO 400

ISO 600



Test Idea and Hypothesis

Hybrid:

- No color bleeding and leakage, particularly through the center
- Performs better when the given image features captured in our fixed features \mathbf{z}

Baseline CNN:

- Perform better in more types of images (Lower validation loss)
- Color leakage, breaking boundary



Datasets Used



CIFAR dataset: Large number of samples (60k images), but low quality (32x32)
Hard for models to learn features



Oxford flower dataset: Small dataset size (~10k images), high resolution (128x128)
Dataset of only UK flowers, not generalizable



STL10 dataset: Dataset inspired by CIFAR, higher resolution (96x96)
Large number of samples: 100k unlabeled images,
5k labeled training images
Good dataset with complex images, inconsistent colorization of images



Data Preprocessing

1. Conversion to RGB

The images grayscale part are extracted using broadcast standards rgb weights ; then, colorization models are applied to analyze their performance.

$$\mathbf{w} = [0.299 \ 0.587 \ 0.114]^\top$$

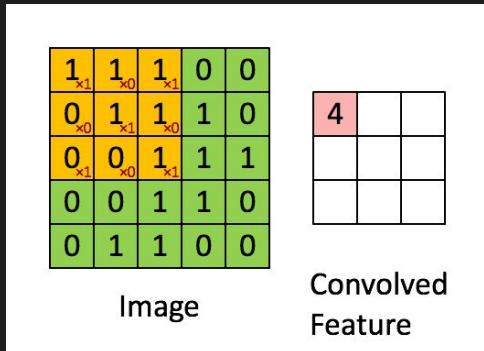
2. Resizing to fit tensor sizes: resize to 144×144 , then center-crop to 128×128 .

3. For better training:

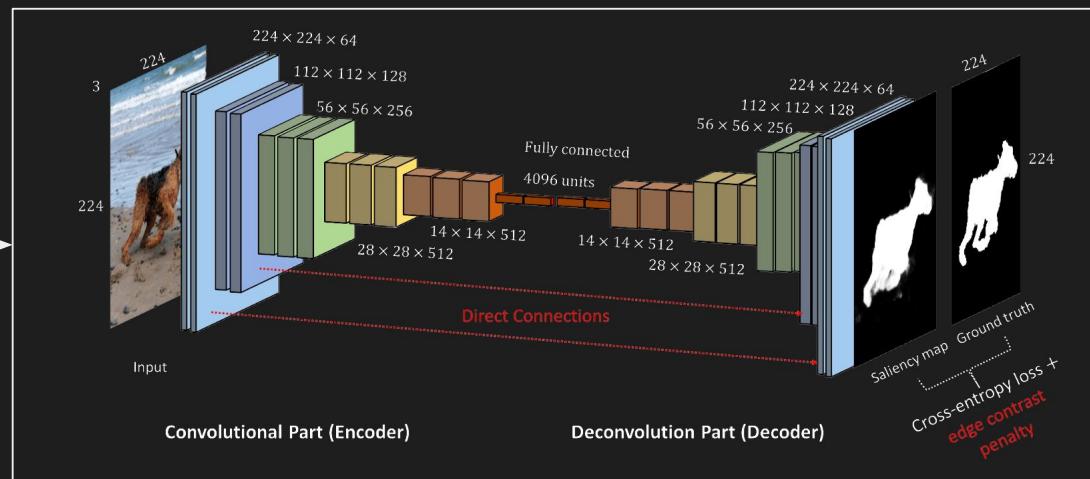
- Add random horizontal flips.
- Add small random rotations (± 10 degrees).

CNN's and Encoder-Decoder Arch.

- CNNs can have Encoder-Decoder structures (somewhat like autoencoders)
- **Convolution:** Large features → compressed representation
 - Learn important parts of an image
- **Deconvolution:** Recreate image



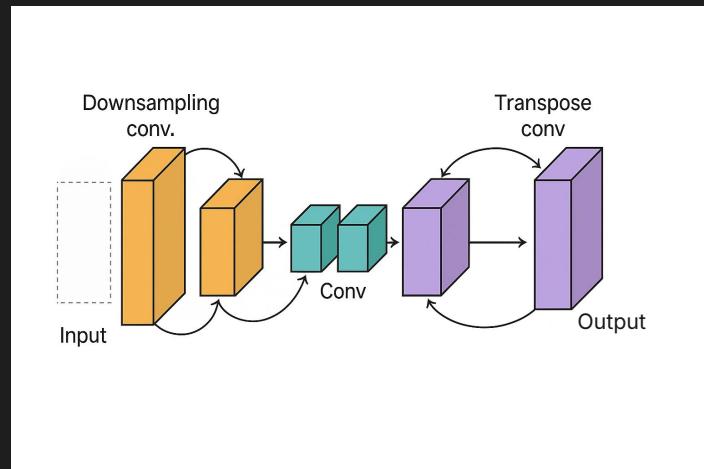
Kernel Demonstration



Generalized CNN with image rebuilding

CNN Encoder-Decoder Arch.

- Input: luminance L
- Encoder: 3x Conv(4x4, stride 2): 1 \rightarrow 32 \rightarrow 64 \rightarrow 128
- Bottleneck: 2x Conv(3x3): 128 \rightarrow 128 \rightarrow 128
- Decoder: 3x ConvTranspose(4x4, stride 2): 128 \rightarrow 64 \rightarrow 32 \rightarrow 2
- Output: predicted chroma coefficients (α, β)





Optimizer

- Our model uses the Adam (Adaptive Moment Estimation) algorithm for optimizing the parameters
- Adam: 2 considerations
 - Momentum optimization (gradient descent that considers the intensity of previous gradients)
 - RMSProp optimization (dynamic learning rate using previous gradients with an emphasis on recent iterations)
- Combining the 2 results in:
 - Adaptable learning rate
 - More efficient optimization
 - Quicker gradient descent

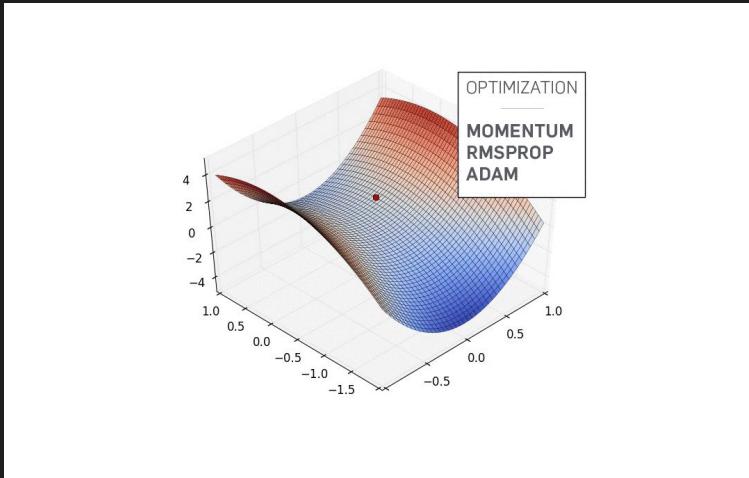


Diagram courtesy of Ayoosh Kathuria
<https://www.digitalocean.com/community/tutorials/intro-to-optimization-momentum-rmsprop-adam>

Baseline CNN Formulation

Input:

$$\mathbf{h}^{(0)} = Y \in \mathbb{R}^{H \times W \times 1}$$

Features Extraction and Normalization:

For layer $\ell = 1, \dots, L - 1$

$$z_k^{(\ell)} = \text{BN} \left(\langle K^{(\ell,k)}, \mathbf{h}^{(\ell-1)} \rangle_F + b_k^{(\ell)} \right)$$

Activation function: Linear \rightarrow Non-Linear.

For layer $\ell = 1, \dots, L - 1$

$$h_k^{(\ell)} = \text{ReLU}(z_k^{(\ell)}) = \begin{cases} 0, & z_k^{(\ell)} \leq 0, \\ z_k^{(\ell)}, & z_k^{(\ell)} > 0. \end{cases}$$

Loss
Function(MSE)

$$\mathcal{L} = \frac{1}{N} \sum_{i,j} \| \mathbf{v}_{ij} - \hat{\mathbf{v}}_{ij} \|^2$$

Output Layer:

$$\mathbf{z}^{(L)} = \begin{bmatrix} \mathbf{z}_1^{(L)} \\ \mathbf{z}_2^{(L)} \end{bmatrix} = \begin{bmatrix} \langle K^{(L,1)}, \mathbf{h}^{(7)} \rangle_F + b_1^{(L)} \\ \langle K^{(L,2)}, \mathbf{h}^{(7)} \rangle_F + b_2^{(L)} \end{bmatrix}$$

Parameter Extraction

$$\begin{bmatrix} \boldsymbol{\alpha} \\ \boldsymbol{\beta} \end{bmatrix} = \begin{bmatrix} \mathbf{z}_1^{(L)} \\ \mathbf{z}_2^{(L)} \end{bmatrix}$$

Model

$$\hat{\mathbf{v}} = \frac{Y}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} + \boldsymbol{\alpha} \mathbf{u}_1 + \boldsymbol{\beta} \mathbf{u}_2$$

Hybrid Formulation

Input:

$$\mathbf{h}^{(0)} = Y \in \mathbb{R}^{H \times W \times 1}$$

Features Extraction and Normalization (hidden layer): **For layer $\ell = 1, \dots, L - 1$**

$$h_k^{(\ell)} = \text{ReLU} \left(\text{BN} \left(\langle K^{(\ell,k)}, \mathbf{h}^{(\ell-1)} \rangle_F + b_k^{(\ell)} \right) \right) \quad \mathbf{z}_{\text{fixed}}^{(m)} = \langle K_{\text{fixed}}^{(m)}, Y \rangle_F$$

Output Layer: *predicting a,b*

$$\mathbf{H}^{(L)} = \begin{bmatrix} \mathbf{a} \\ \mathbf{b} \end{bmatrix} = \langle K^{(L)}, \mathbf{h}^{(L-1)} \rangle_F + b^{(L)}$$

Model:

$$\hat{\mathbf{v}} = \frac{Y}{\mathbf{w}^\top \mathbf{w}} \mathbf{w} + (\mathbf{a} \cdot \mathbf{z}_{\text{fixed}}) \mathbf{u}_1 + (\mathbf{b} \cdot \mathbf{z}_{\text{fixed}}) \mathbf{u}_2$$

Loss
Function(MSE):

$$\mathcal{L} = \frac{1}{N} \sum_{i,j} \|\mathbf{v}_{ij} - \hat{\mathbf{v}}_{ij}\|^2$$

Kernels:

- Laplacian: Clusters
- Sobel X: vertical lines
- Sobel Y: horizontal lines



Baseline CNN Model

Pros

- Conditional logic (continuous function):
 - Can learn non-linear relationships.
 - Non-linear mapping from Luma to chroma
- Generalized Features
- Adaptive Kernels

Cons

- Misalignment and Gradient leak
- Unconstrained
 - Entire search space is 2D plane
 - No Structural assumptions
 - Learn structure and color
 - Often results in desaturation under MSE



Hybrid Model

Pros

- Lower (MSE) error during training
- Pre-defined feature extraction layer
- Edge Alignment (No gradient leak)
- Tighter constraint
 - Search space is defined by **z**
 - Only need to learn the coefficients

Cons

- Longer training time (more parameters)
- Linearity
- Can't do conditional logic(continuous linear map)
- Noise sensitivity



04 Results

AUTO

ISO 90

ISO 100

ISO 200

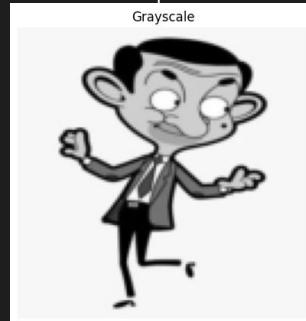
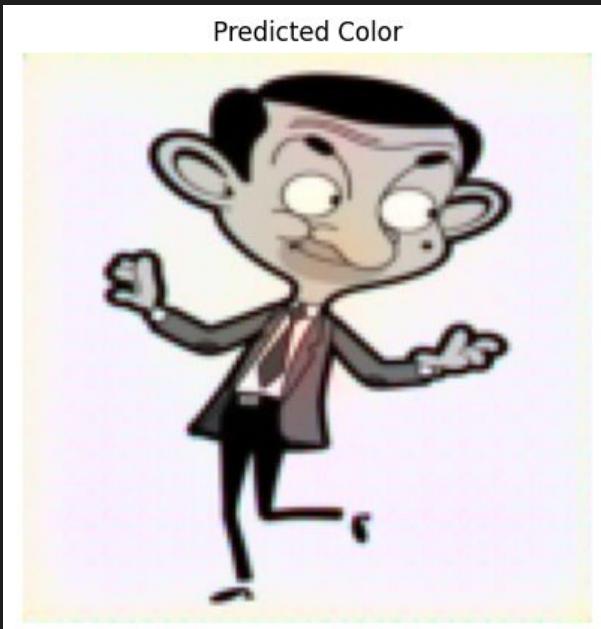
ISO 400

ISO 600

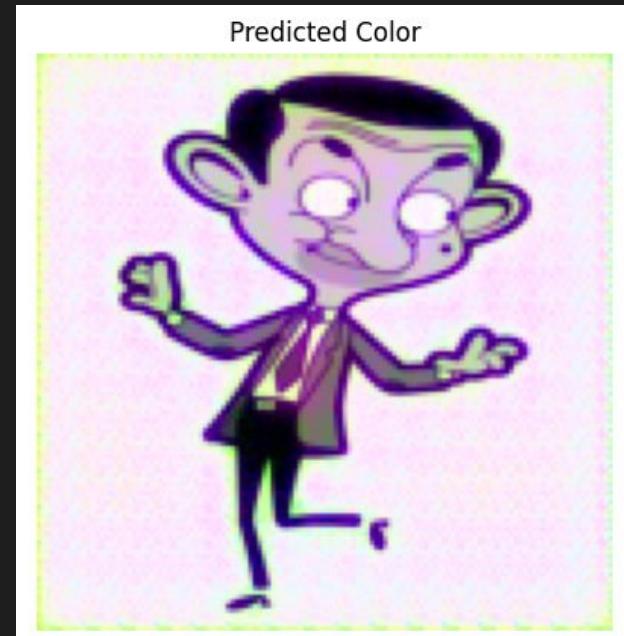
Hybrid (SLR + CNN) vs Base (CNN) Models

Gradient Leakage!

Hybrid Model



Base Model



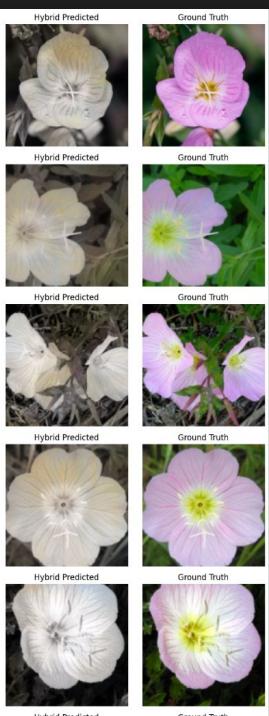
Training on CIFAR10

Hybrid Model

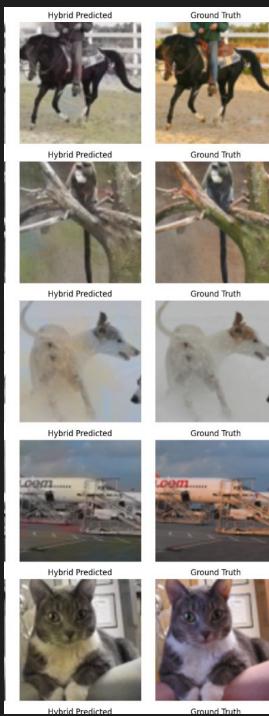
CIFAR10



Flowers



STL10



Base Model

Flowers



STL10



Training on Flowers

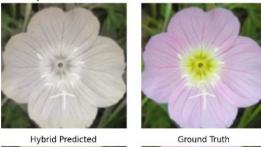
Gradient Leakage!

Hybrid Model

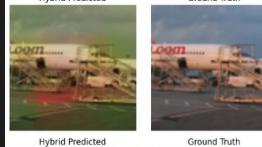
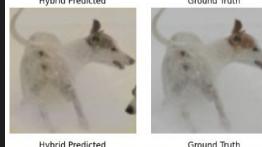
CIFAR10



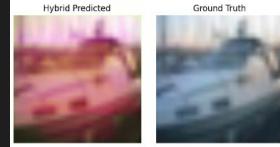
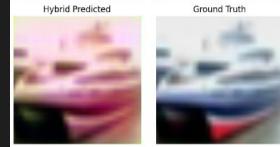
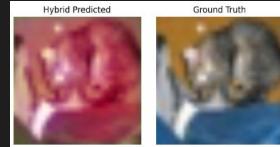
Flowers



STL10



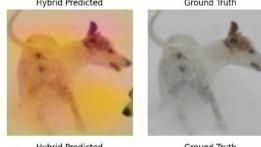
CIFAR10



Flowers



STL10



Base Model



Best Performing

Hybrid Model

Training on STL10

Base Model

CIFAR10

Flowers

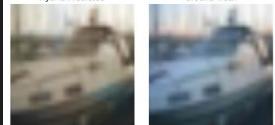
STL10



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth



Ground Truth

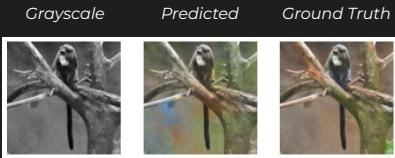


Ground Truth

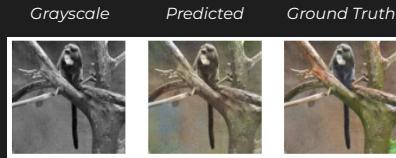
Effect of Data Augmentation

Hybrid Model

Without



With Data Aug

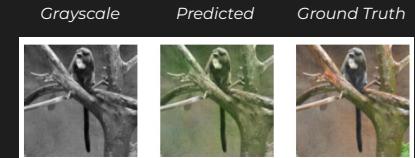


Base Model

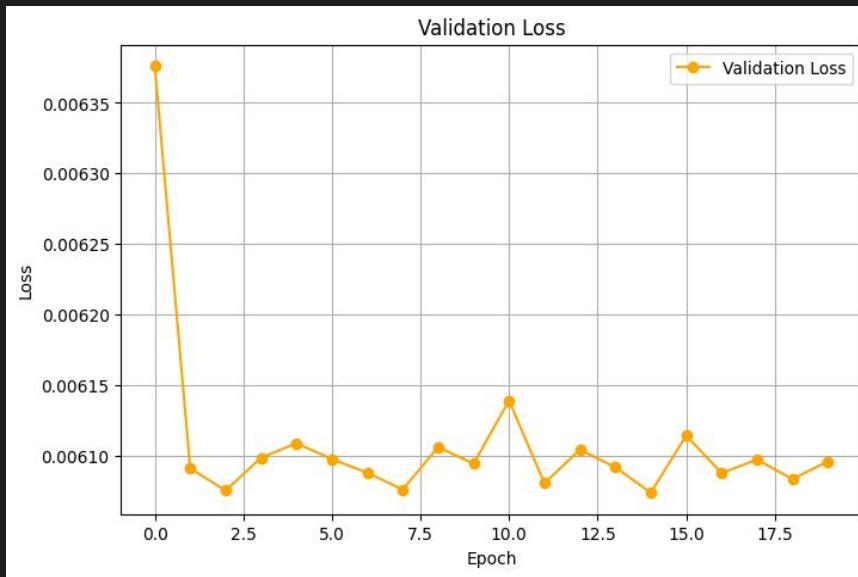
Without



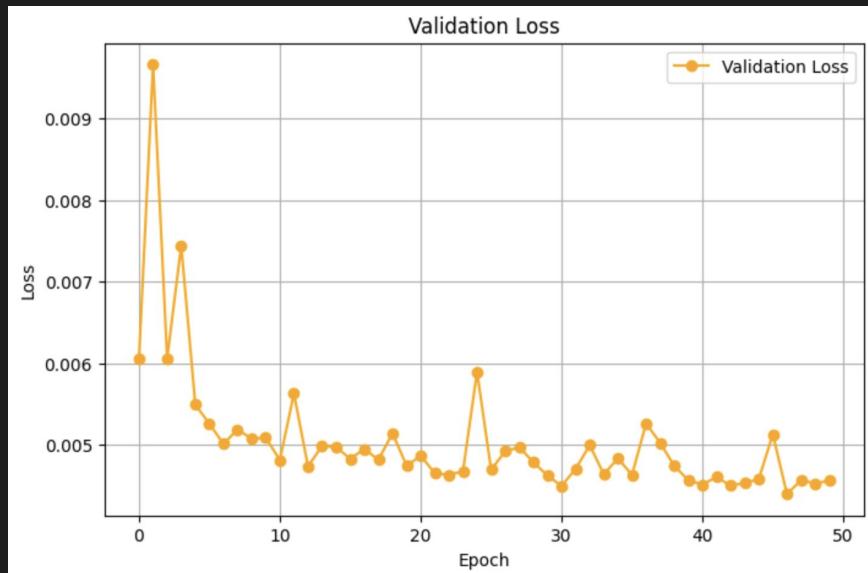
With Data Aug



Effect of no. layers on performance of CNN



1 layers CNN with hybrid structure
Underfitting



8 layers CNN with bottleneck layers,
computed at 128 x 128
Validation loss converges



05 Discussion

AUTO

ISO 90

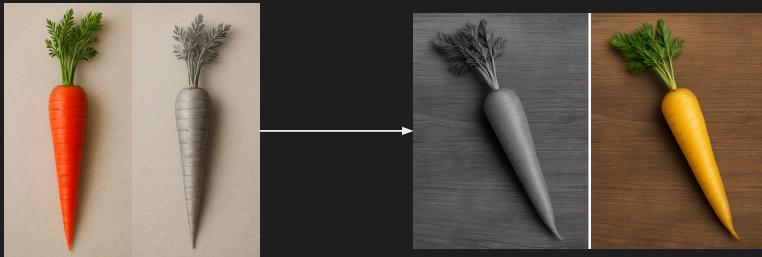
ISO 100

ISO 200

ISO 400

ISO 600

General Issues with Datasets



We relied on public datasets for tasks like detection and classification, not image colorization. This introduced a few notable issues:

- The colors of objects were not consistent
- Images were not large enough for our model to learn specific color shades
- We made grayscale images from RGB ones, but grayscale ≠ corrected monochrome images

There is a need to craft a proper benchmarking dataset that:

1. Ensures standardized colors across objects
2. Contains images of sufficient resolution
3. Has standalone RGB and grayscale images for each item

Specific Issues with Datasets

- CIFAR10:
 - Required extensive upsampling (bilinear interpolation)
- Oxford Flowers:
 - Small amount of samples led to overfitting
- STL10:
 - Large amount of samples led to inconsistency amongst objects within training images

Possible Improvements

- Ensemble learning
 - Combine multiple models together
- Test with a different module
 - Generative Adversarial Network to predict colors differently
- Have a more robust Dataset
 - Small datasets lead to overfitting
 - More high resolution images for more accurate testing and predictions
- Use Lab Space
- Different Loss Functions
- Classification
- Apply post-processing to remove artefacts using
 - Joint bilateral filtering
 - Semantic labels

Our Takeaways

- The **SLR + CNN model** is effective at using familiar local patterns and structural features
- Makes “average” guesses
- Confirmed our hypothesis that the **SLR + CNN model** is a potential solution for the gradient leakage problem (no color bleeding)
- Both models have limitations when it comes to prediction accuracy, but show potential



06 Conclusion

AUTO

ISO 90

ISO 100

ISO 200

ISO 400

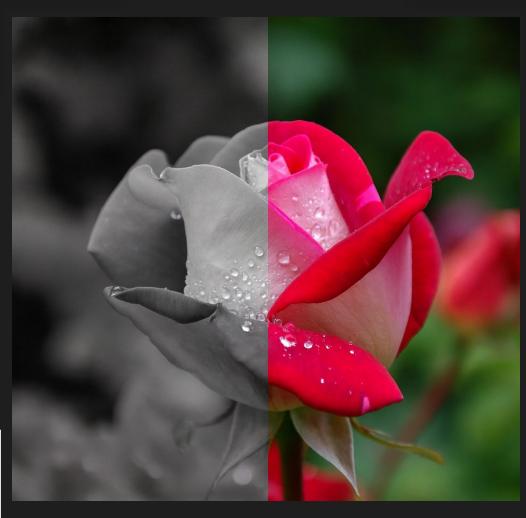
ISO 600

Conclusion

- The use of automated methods for image colorization can help make this task more convenient and accessible

Our hybrid SLR + CNN model shows potential for image colorization with its improvements in identifying key features within images and predicting RGB values from grayscale inputs, addressing the color drift and saturation issues within the baseline model.

- Our findings emphasize the value in advancing image colorization methods to streamline not only archival and restoration practices, but also assist creative professionals in enhancing their work



AUTO

ISO 90

ISO 100

ISO 200

ISO 400

ISO 600

Demo

<http://localhost:5000/>





Thank You

Questions?



AUTO

ISO 90

ISO 100

ISO 200

ISO 400

ISO 600